



University of the Philippines

Single-Source Shortest Paths: From Dijkstra to DMMSY

A CMSC 142 portfolio study comparing five SSSP algorithms in pure-Python stdlib, with DMMSY (2025) as the centerpiece.

Zildjian E. California¹ Rey Marvin C. Rizal¹

¹ Department of Mathematics, Physics, and Computer Science,
College of Science and Mathematics, University of the Philippines - Mindanao

CMSC 142 Portfolio Presentation – Assoc. Prof. Cinmayii G. Manliguez,
Ph.D. · May 18, 2026

- 1 Problem and Motivation
- 2 Baseline Algorithms
- 3 DMMSY (Centerpiece)
- 4 Experimental Evaluation
- 5 Reflection and Author Contributions

Problem and Motivation

Problem

Given a weighted graph $G = (V, E)$ with edge weights $w: E \rightarrow \mathbb{R}$ and a source vertex $s \in V$, compute

$$d(v) = \min_{\pi: s \rightarrow v} \sum_{e \in \pi} w(e) \quad \text{for every } v \in V.$$

- Edge orientation: directed for Dijkstra, Bellman–Ford, A^* , DMMSY; undirected for Thorup.
- Weight class: non-negative reals for Dijkstra / A^* / DMMSY; signed reals for Bellman–Ford; non-negative integers for Thorup.
- Output: a distance map $d: V \rightarrow \mathbb{R}$ (and, optionally, a predecessor map for path reconstruction).

- **Dijkstra (1959) [1]** – the foundational non-negative-weight algorithm; binary-heap relaxation.
- **Bellman–Ford (1958) [2]**, drawing on Ford (1956) [3] – handles negative weights; detects negative cycles.
- **A* (1968) [4]** – heuristic-guided search; degenerates to Dijkstra at $h \equiv 0$.
- **Thorup (1999) [5]** – linear-time undirected SSSP for non-negative integer weights in the word-RAM model.
- **DMMSY (2025) [6]** – deterministic $O(m \log^{2/3} n)$ directed SSSP; **the centerpiece of this study.**

We re-implement all five in pure-stdlib CPython 3.11+ and benchmark each against Dijkstra on a documented suite.

Baseline Algorithms

Problem class

Non-negative weighted directed (or undirected) graphs; single source s .

Design idea

Maintain a frontier set B in a min-heap keyed by tentative distance; repeatedly extract the closest unsettled vertex, settle it, and relax its outgoing edges.

Complexity

Time $O((n+m) \log n)$; space $O(n)$.

Implementation

`src/sssp/dijkstra.py`;
heapq-backed; 7 tests, 100% line coverage. Source paper: [1].

Problem class

Directed graphs with possibly negative edge weights; detect negative-weight cycles reachable from s .

Design idea

Repeat full-graph edge relaxation $n - 1$ times; one extra pass diagnoses any reachable negative cycle.

Complexity

Time $O(nm)$; space $O(n)$.

Implementation

```
src/sssp/bellman_ford.py;  
reports  
NegativeCycleReport; 100%  
line coverage. Source papers: [2],  
[3].
```


Problem class

Goal-directed shortest path with a domain-specific heuristic $h(v)$ lower-bounding $d(v, \text{goal})$.

Design idea

Order the frontier by $f(v) = g(v) + h(v)$ instead of $g(v)$ alone; admissible h preserves optimality.

Complexity

Time $O((n + m) \log n)$ worst case; expands fewer nodes than Dijkstra when h is informative.

Implementation

```
src/sssp/astar.py,  
src/sssp/heuristics.py;  
ships    manhattan,    zero;  
 $A^*(h \equiv 0) \equiv \text{Dijkstra}$  tested.  
Source paper: [4].
```

Problem class

Undirected graphs with non-negative integer edge weights; word-RAM model.

Design idea

Hierarchical bucket structure exploiting integer weight bits and constant-time word operations to avoid the sorting bottleneck of comparison-based SSSP.

Complexity

Time $O(n + m)$ in the word-RAM model.

Implementation

`src/sssp/thorup99.py`; 11 tests, 99% line coverage. *Python* $\neq$ *word-RAM*: documented in ADR-002. Source paper: [5].

DMMSY (Centerpiece)

Setting

Directed SSSP with non-negative real weights; labels use only comparison and addition.

Barrier

Dijkstra-style methods still pay for global priority ordering, giving the sparse-graph comparison baseline $O(m + n \log n)$.

DMMSY claim

Duan–Mao–Mao–Shu–Yin break the barrier deterministically:

$$O\left(m \log^{2/3} n\right).$$

Portfolio role

Centerpiece algorithm. Source paper: [6]. Implementation: `src/sssp/dmmsy/`.



- **Transform:** reduce high-degree vertices so work charges uniformly per arc.
- **FindPivots:** shrink the frontier to useful recursion roots.
- **BMSSP + BlockList:** solve bounded subproblems without globally sorting every label.

Contract

$\text{BMSSP}(G, \ell, B, S, d, k, t)$ solves labels below boundary B from source set S , then returns a smaller boundary B' and settled set U .

Level 0

BASECASE: bounded Dijkstra from one source, stopping after $k + 1$ discoveries or boundary exhaustion.

Level > 0

- 1 FINDPIVOTS shrinks S to pivots P .
- 2 BlockList pulls bounded batches below B .
- 3 Recursive calls settle batches and relax outgoing edges.

Key idea: local bounded batches replace one global sorted frontier.

Paper parameterization

$$k = \left\lfloor (\log_2 n)^{1/3} \right\rfloor$$

$$t = \left\lfloor (\log_2 n)^{2/3} \right\rfloor, \quad \ell_{\max} = \left\lceil \frac{\log_2 n}{t} \right\rceil$$

Balances pivot work, recursion depth, and block-list operations.

Asymptotic result

Deterministic directed SSSP:

$$O\left(m \log^{2/3} n\right) \text{ time}$$

Space stays $O(n + m)$.

Python implementation boundary

CPython recursion overhead is real. Default large numeric benchmark runs use the documented fast path; `Weight` inputs and explicit parameters exercise recursive BMSSP.


Experimental Evaluation

Dataset

Seed 2026; 12 benchmark cases; geometric and Barabasi–Albert families at $n \in \{64, 256, 1024\}$ with sparse/dense profiles.

Timing

60 CSV rows; setup and file IO excluded; one recorded repeat per case-algorithm pair; Dijkstra is the normalization baseline.



Median runtime ratio by size bucket

Algorithm	Small $n = 64$	Medium $n = 256$	Large $n = 1024$
Dijkstra	1.00×	1.00×	1.00×
Bellman–Ford	1.92×	3.48×	3.87×
A*	0.80×	0.29×	0.50×
Thorup (1999)	1.84×	1.61×	1.36×
DMMSY (2025)	11.28×	5.87×	1.02×

- A* benefits from an informative geometric heuristic; Bellman–Ford pays for repeated edge sweeps.
- Thorup stays above Dijkstra because CPython is not the word-RAM model.
- DMMSY clears the T-17 sanity gate at large size: max DMMSY/Dijkstra ratio is 1.093x for $n \geq 1000$.

Reflection and Author Contributions

What the comparison teaches

No SSSP algorithm dominates every setting; the right choice depends on weight signs, direction, integer assumptions, heuristic quality, and runtime model.

Honest DMMSY framing

The centerpiece is the deterministic $O(m \log^{2/3} n)$ comparison-addition bound. CPython does not show a recursive BMSSP wall-clock win; large numeric runs use the ADR-003 fast path to keep NFR-F green.

AI assistance disclosure

AI tools assisted with drafting, tests, artifacts, and prose. The source papers, validation evidence, and final judgments remain the human authors' responsibility.

Zildjian E. California

- Repository foundation, governance, graph/weight/IO/generator base, Dijkstra, A*, DMMSY, benchmarks, report, and deck authorship.
- DMMSY centerpiece: transform, block list, pivots, BMSSP, driver, ADR-003/ADR-004, and presentation opening/centerpiece.
- Current commit history: authored the foundation, Dijkstra/A*, DMMSY stack, experimental-evaluation artifacts, report manuscript, and deck scaffold.

Rey Marvin C. Rizal

- Bellman–Ford implementation with negative-cycle reporting; feature issue #4 closed as completed.
- Thorup 1999 runtime-model work and hierarchical-bucket SSSP contribution; feature issue #5 closed as completed.
- Presentation ownership: Bellman–Ford, Thorup, and experimental-results delivery; joint Q&A with California.

Contribution split follows the provided commit history and the closed Rizal-assigned issues #4 and #5.

Honor. Excellence. Service.

Questions?

Zildjian E. California zecalifornia.com

Rey Marvin C. Rizal

Repository:

<https://github.com/zecalifornia-ph/sssp-modern>



University of the Philippines

The five source papers cited above are listed below; the Beamer class [7] and the UP Visual Identity Guidebook [8] are credited for the deck infrastructure.

- [1] **E. W. Dijkstra**. “A Note on Two Problems in Connexion with Graphs”. In: *Numerische Mathematik* 1.1 (1959), pp. 269–271. DOI: [10.1007/BF01386390](https://doi.org/10.1007/BF01386390).
- [2] **Richard Bellman**. “On a Routing Problem”. In: *Quarterly of Applied Mathematics* 16.1 (1958), pp. 87–90. DOI: [10.1090/qam/102435](https://doi.org/10.1090/qam/102435).
- [3] **L. R. Ford**. “Network Flow Theory”. In: P-923 (1956). Also available at <https://www.rand.org/pubs/papers/P923.html>. URL: <https://apps.dtic.mil/sti/tr/pdf/AD0422842.pdf>.
- [4] **Peter E. Hart, Nils J. Nilsson, and Bertram Raphael**. “A Formal Basis for the Heuristic Determination of Minimum Cost Paths”. In: *IEEE Transactions on Systems Science and Cybernetics* 4.2 (1968), pp. 100–107. DOI: [10.1109/TSSC.1968.300136](https://doi.org/10.1109/TSSC.1968.300136).

- [5] **Mikkel Thorup**. “Undirected Single-Source Shortest Paths with Positive Integer Weights in Linear Time”. In: *Journal of the ACM* 46.3 (1999), pp. 362–394. DOI: [10.1145/316542.316548](https://doi.org/10.1145/316542.316548).
- [6] **Ran Duan et al.** “Breaking the Sorting Barrier for Directed Single-Source Shortest Paths”. In: (2025), pp. 36–44. DOI: [10.1145/3717823.3718179](https://doi.org/10.1145/3717823.3718179).
- [7] **Till Tantau**. *The Beamer Class User Guide*. 2004.
- [8] **University of the Philippines**. *UP Visual Identity Guide 2017*. Accessed on April 2, 2026. URL: <https://up.edu.ph/wp-content/uploads/2017/03/VIG-layout-Jan-2017B-compressed.pdf>.

Artifacts to cite during Q&A

- CSV: `bench/results/sssp-benchmark-seed-2026.csv`
- Figure: `bench/figures/runtime-by-algorithm.{pdf,png}`
- Tests: 147-test suite passed after benchmark plot generation; T-17 records max DMMSY/Dijkstra ratio $1.093\times$ for $n \geq 1000$.
- Report anchors: Table `tab:results-by-class` and Figure `fig:runtime-by-algorithm`.

If asked why DMMSY is still the centerpiece

The portfolio centers the algorithmic result and implementation structure, not a CPython timing victory: correctness and comparison-addition fidelity are validated separately from the fast-path benchmark policy.